

Product Specification

What the product does, how each part works, and what we are deliberately leaving out. Written to be read start to finish in about twenty minutes.

AUDIENCE

Founding team

STATUS

Working draft

FIRST EDITION

Recruiting

UPDATED

29 August 2026

CONTENTS

- | | |
|-----------------------------|----------------------------------|
| 1 What Sourcewell is | 6 One engine, many sectors |
| 2 The words we use | 7 What we deliberately do not do |
| 3 The journey, end to end | 8 Open decisions |
| 4 The product, area by area | 9 Where we are today |
| 5 The autonomy dial | |

1 What Sourcewell is

Sourcewell finds people worth contacting, writes to each of them personally, follows up over several days, and brings every reply back to one place. A recruiter describes the role they are filling. The system searches licensed people databases, ranks everyone it finds against that description, drafts a tailored first message for each person, sends it from the recruiter's own mailbox once a human approves it, follows up on a schedule, and hands the conversation back the moment someone shows real interest.

The work it replaces is repetitive and slow: searching, reading profiles, writing a note that sounds like a person wrote it, remembering to follow up, and tracking who answered. Done properly that is several minutes per candidate, a few hundred times per role. Sourcewell runs that loop continuously in the background so the team reviews work instead of producing it.

Who it is for

Talent teams at large, regulated employers (banks, insurers, and similar), and staffing agencies working several clients at once. The person who buys it leads talent acquisition. The people who can block the purchase are security, compliance, and legal, which is why approval gates, audit records, and sending limits are part of the core product rather than an enterprise upsell we add later.

What makes it different

Most tools in this space do one slice: a database to search, a sequencer to send, an inbox to read. The team stitches them together and carries the judgment between them. Sourcewell ties a single target (one open role) to a pipeline that stays filled, ranked, contacted, and answered, as one supervised system with one record of what happened and why.

The second difference is posture. The system is built to be supervised, not trusted blindly. Nothing reaches a real person without a human having approved it, until the team explicitly decides otherwise, campaign by campaign.

2 The words we use

These are the objects the product is built from. They appear in the interface with these names, so it is worth agreeing on them early.

TERM	WHAT IT MEANS
Person	Someone who signs in. One login works everywhere they work, so a recruiter at two agencies has one account rather than two.
Organization	The customer. Billing, the do-not-contact list, data-provider contracts, and the audit record live here. Most customers never think about it; it is simply their company.
Workspace	The isolation boundary. Contacts, campaigns, and conversations live inside one and never cross to another. Shown as Client, Department, or Team depending on how the customer works.
Contact	A person we might reach out to. Called a candidate in the recruiting edition.
Campaign	One outreach effort: who we want, what we send, and the rules we send under. In recruiting, one open role.
Enrollment	One contact being worked inside one campaign, carrying its own state, score, and history.
Sequence	The ordered messages a campaign sends, each with a channel and a delay.
Seat	A person's own connected sending account: their work mailbox or their LinkedIn account. It belongs to the person, and each campaign names the seat it sends from.
Fit score	How well a contact matches a campaign, from 0 to 100, always with a short written reason.
Suppression	The do-not-contact list. Company-wide by default, and permanent once someone is on it.
Partner	A company that puts their own brand on the product and sells it to their own customers. A partner sets defaults; they do not see their customers' pipelines.

3 The journey, end to end

This is one pass through the whole product, in the order a customer experiences it. Each step is described in full in section 4.

An administrator signs in and invites the team. Each recruiter connects their work mailbox, and optionally their own LinkedIn account, because messages send from the person, not from us. The team sets its ground rules once: who must never be contacted, how many messages a day is acceptable, what hours are acceptable, and how the writing should sound.

A recruiter starts a campaign by pasting a job description, pulling a posting from LinkedIn, or just describing the role in a sentence. The system turns that into a structured description of who we are looking for, which the recruiter reviews and corrects. This step matters more than it looks, because everything downstream is measured against it.

The recruiter then confirms the audience, picks which sources to search, builds the message sequence, chooses which account sends, and sets the autonomy level. From that point the campaign runs on its own clock.

The system searches continuously, brings back people who match, scores each one against the role with a written reason, and drops anyone already in the pipeline or on the do-not-contact list. The recruiter reviews a ranked shortlist and approves who to pursue. For each approved person, the system drafts a first message and holds it for approval. Approved messages send from the recruiter's own account inside the sending limits, and follow-ups go out on schedule until the person replies or opts out.

Replies from any channel land in one inbox, grouped as a single conversation per person. The system reads each reply, decides whether it is interest, a question, a decline, or an opt-out, pauses the follow-ups, and drafts a response. When someone is genuinely interested, the conversation is handed to the recruiter with a summary of what was said, and the automation stops for that person.

4 The product, area by area

4.1 Accounts, onboarding, and connections

A customer signs up as an organization and gets one workspace to start with. People are invited by email, and one login works across every company a person belongs to, so a contractor working for two agencies signs in once and switches between them.

Access has two parts. An organization role covers the company: an admin manages billing, members, and company-wide rules; a member works the pipeline; a compliance role reads the audit record without touching campaigns. On top of that, people are granted access to particular workspaces. Organization admins reach every workspace by default, and everyone else sees only what they have been granted.

Each person connects their own sending accounts. The work mailbox is the important one, because email is the channel that scales and stays defensible, and a message from a real recruiter's address arrives better than one from a shared marketing domain. LinkedIn is optional and connected per person, using their own account. First-run onboarding walks a new workspace through three things: connect an account, add or find some people, create a first campaign.

4.2 Finding people

There are three ways people enter a workspace. The team can import a list they already have. A recruiter can search interactively, typing what they want in plain language, adjusting the filters, and picking who to keep. Or a running campaign can search on its own schedule and keep adding people who match.

Searches run against licensed people databases and, where a recruiter has connected their own LinkedIn account, against LinkedIn itself. Several sources can be searched at once, results are merged so the same person never appears twice, and each result carries a note of where it came from. We do not keep a copy of these databases. We search them live and store only the people a customer chooses to keep, which is both cheaper and a much easier privacy position to defend.

Before anything is kept, we check contact details, drop anyone on the do-not-contact list, and drop anyone already in the pipeline. A campaign keeps topping itself up as long as it is active, so the recruiter sees new candidates arriving rather than a list that goes stale.

4.3 Ranking and fit

Every person is scored from 0 to 100 against the campaign's description, and every score comes with a sentence explaining it in plain language, citing the person's actual background. The score weighs the things a recruiter would weigh: title, skills, seniority, function, employer, industry, company size, and location. Exclusions are absolute: if a campaign says never contact anyone at a named competitor, a match there is disqualified regardless of how good the rest looks.

Whether we can actually reach someone is tracked separately from how well they fit, because they are different problems. A perfect match with no verified email is still a perfect match; it just needs more work to reach. Mixing the two into one number hides that.

The same scoring runs when a recruiter is still editing the audience, so the live count of matching people they see while typing is the same count the campaign will act on. That consistency is a deliberate constraint and worth protecting.

4.4 The campaign

The campaign is the control surface. Everything the system does for a role is a consequence of what is set here, so the settings screen is effectively the product's policy screen.

SETTING	WHAT IT CONTROLS
Goal	The role in plain words. Drives how the assistant writes and reasons.
Audience	Who we are looking for. Drives both searching and ranking.
Sources	Which databases and channels this campaign is allowed to search.
Pipeline target	How many live candidates to keep. Tells the system when to search more and when to stop.
Score threshold	How good a match has to be before anyone proposes it.
Sequence	The messages, their order, their channel, and the wait between them.
Sending account	Whose mailbox or LinkedIn account these messages come from.
Autonomy	Which approvals a human still has to give. See section 5.
Limits	Daily send caps, acceptable hours, and any campaign-specific exclusions.

A campaign moves through four states. It starts as a **draft** that costs nothing and sends nothing. Launching makes it **active**, and the launch step is the real review gate, the moment a person confirms the whole configuration rather than each setting individually. **Pausing** stops all new work immediately while holding everything in flight, so it doubles as the emergency stop. **Completing** retires it while keeping the history.

4.5 Messages and sequences

A sequence is an ordered set of touchpoints. Each one names a channel, a wait before it goes out, and either a template or an instruction to the assistant. A typical recruiting sequence is an email on day zero, a LinkedIn message on day three, and a short final email on day six.

Messages are written per person, not merged from a template. The assistant is given the person's real background, the role, and the team's writing guidance, and is asked for something short and specific. Teams can also upload approved templates and pin a sequence step to one, which matters for organizations whose

legal team wants to see the wording in advance.

Every generated message is checked before it can queue: length and format, whether it claims anything about the person that is not in their profile, whether it uses banned language, whether it matches the team's voice, whether it carries the required sender identification and unsubscribe wording, and whether we have said something too similar to this person before. A message that fails a check does not reach the approval queue.

4.6 Approvals

Approvals are where a person decides what the system is allowed to do, and there are two of them. The first covers who we pursue: the agent proposes a ranked list of people, each with a fit score and a short explanation, and the recruiter accepts or dismisses them individually or in bulk. The second covers what gets sent: every message the agent writes waits as a draft until someone approves, edits, or rejects it. Both gates can be relaxed per campaign as a team grows comfortable, but a new campaign starts with both on, and nothing reaches a real person without someone having seen it.

Dismissing a candidate and blocking a person are separate actions, deliberately. Saying someone is wrong for this role means only that; it does not follow them to other roles or other clients. Adding someone to the do-not-contact list is a distinct, heavier decision that applies everywhere in the organization and does not expire.

4.7 Sending and pacing

Approved messages do not send immediately. They pass through a set of limits designed to keep both the customer's email domain and their team's LinkedIn accounts healthy.

- A daily cap per channel, per sending account, set by the workspace.
- Acceptable sending hours and days, so nothing goes out at 3am on a Sunday.
- A warmup ramp for new accounts, which starts low and increases over the first weeks.
- Small natural variation in timing, so sending does not look mechanical.
- A shared budget when one person works several workspaces, split across them rather than multiplied.

Every campaign names the seat it sends from. If that seat is missing or unhealthy we stop and say so rather than quietly sending from a colleague's account, because a candidate receiving a note from a recruiter they have never heard of is worse than a delayed send.

If a LinkedIn account shows any sign of trouble, we stop using it, tell its owner, and ask them to reconnect rather than pushing through. If a campaign wants LinkedIn but the seat is missing or unhealthy, we fall back to email when we have an address and surface a visible failure when we do not. We never record a message as sent when it was not.

4.8 The inbox and replies

Replies arrive in one inbox, grouped into a single conversation per person even when the thread moves between email and LinkedIn. The inbox also holds drafts waiting for approval, so a recruiter has one queue to work rather than two.

Each reply is read and classified: interested, a question, not right now, or an opt-out. Follow-ups stop the moment anyone replies, so nobody receives a scheduled nudge after they have already answered. Opt-outs are honored immediately and permanently, without waiting for a human. For everything else the assistant drafts a response the recruiter can send, edit, or discard.

Alongside each conversation sits the context a recruiter needs to answer well: who this person is, what we said before, what they scored and why, and which campaign they belong to.

4.9 Handoff

When someone is genuinely interested, the system stops and gives the conversation to a person, with a short summary of what was discussed. Handoff also triggers on its own whenever the assistant is out of its depth: a question about compensation or terms, a topic outside the role, an escalation or complaint, or plain uncertainty. Those handoffs happen regardless of the autonomy setting, because the cost of a confident wrong answer to a real candidate is high and the cost of asking a human is low.

4.10 The assistant

There is one assistant, seen from three places. It designs a campaign from a brief and keeps proposing improvements as results come in. It sources and scores people continuously. It handles live conversations. To a user it should feel like one colleague working in the background, not like three separate features.

Two rules govern it. The first is that the assistant is never load-bearing: every part of the product works without it, in a simpler and more mechanical form, so an outage or a missing key degrades quality rather than stopping work. The second is ownership. Any part of a campaign a person edits by hand becomes theirs, and the assistant will suggest changes there but never make them. Parts the assistant owns it can change on its own, recording what it changed and why. Handing a section back is a single explicit action.

Everything the assistant does is recorded as a readable trail: what it looked at, what it decided, what it changed. A recruiter can open any campaign and see the last week of its work in plain language, and undo anything it did on its own.

The assistant is also available as chat, in context. Asking about a campaign should be able to return not just a sentence but the actual thing being discussed: a funnel, a shortlist, a proposed audience with an Apply button next to it.

4.11 Companies, workspaces, and people

There are exactly two levels above a person, because the product has to keep two promises at the same time.

The first is isolation. An agency running two competing banks needs a hard guarantee that a candidate sourced for one never surfaces for the other, and an enterprise needs departments walled off from each other. That is what a workspace is. It is a boundary rather than a filter, because a filter can be forgotten and a boundary cannot.

The second is that some things have to be company-wide. If someone asks never to be contacted again, that has to hold for every client the agency works, not only the one that reached them. The same goes for the invoice, the data-provider contract, the audit record, and single sign-on. Those belong to the organization.

One level cannot express both, which is why there are two. Everything else that varies is deliberately kept out of the hierarchy: a person is their own login, a seat belongs to a person, and a reseller is a label on a company rather than a layer above it.

WALLED INSIDE A WORKSPACE	SHARED ACROSS THE COMPANY
Contacts and their history	The do-not-contact list
Campaigns and sequences	The plan, credits, and invoice
Conversations and replies	Data-provider contracts
Fit scores and the reasoning behind them	The audit record
Workspace settings and brand voice	Sign-in and team membership

A sending seat sits outside both. It belongs to the person who connected it, and its daily budget is shared across every workspace they work rather than multiplied by them.

Most customers never see any of this. A single team gets one workspace made for them and never meets the word. Agencies and enterprises add workspaces when they need them.

4.12 Compliance and safety

This is the section that decides enterprise deals, and it is best treated as product surface rather than paperwork.

- A complete record of every decision, draft, send, and reply, showing who or what did it, when, and why, exportable on request.
- A do-not-contact list that is company-wide by default, so one opt-out binds every client, with the option of an entry that applies to a single workspace.
- Fair hiring: candidates are assessed only on job-related factors, protected characteristics are never used as criteria, and the reasoning behind every selection is retained.
- Privacy: people can be removed on request, we keep only what a customer imported, and each customer can choose the region their data lives in.
- Outbound rules: every email identifies the sender and carries a working unsubscribe link.
- Account safety: LinkedIn stays inside human-like limits and pauses on any warning; email-only mode exists for customers whose security policy forbids LinkedIn automation entirely.
- A kill switch that stops one campaign or the entire account immediately.

4.13 Reporting

Reporting answers three questions. Is the pipeline healthy, shown as a funnel from sourced to contacted to replied to handed off, per campaign and per person sending. Which messages actually work, comparing sequences and steps by reply rate. And are the accounts healthy, showing daily volume against caps and flagging any seat in trouble. A notification feed surfaces the things needing attention: new replies, waiting approvals, and completed handoffs.

4.14 Plans and usage

Billing is per organization, on a monthly plan with an included allowance of credits shared across all its workspaces. Three things consume credits: an email sent, a LinkedIn message sent, and a candidate sourced. LinkedIn costs more than email because the underlying capacity is scarcer.

Usage is always visible, both as a running total and broken down by action, and we warn before the limit rather than after. Going over the allowance does not stop work today; it is surfaced and reconciled. Paid plans are self-serve through a standard checkout, and we never handle card details.

5 The autonomy dial

Autonomy is one setting per campaign with three positions. It is the single most important control in the product, because it decides how much a customer is trusting us with, and it lets a nervous buyer start safely and move up as evidence accumulates rather than deciding on day one.

MODE	WHAT THE ASSISTANT DOES	WHAT A PERSON STILL DOES
Manual	Suggests only. Finds and ranks people, drafts messages, proposes changes.	Approves who to pursue and every message, and applies every change by hand.
Assisted	Sources and drafts continuously, and adjusts the parts of the campaign it owns.	Approves every message before it sends.
Autopilot	Sources, drafts, and sends inside the campaign's limits.	Steps in on replies, and on anything the assistant escalates.

FIXED IN EVERY MODE

Opt-outs are honored immediately. The do-not-contact list is enforced before every send. Sending limits and quiet hours always apply. Handoff triggers (compensation, off-topic, escalation, uncertainty) always fire. Every action is recorded. Autopilot buys speed, never a reduction in safety.

New campaigns start in Assisted. Manual exists for regulated teams and for the first week of a pilot, when a customer wants to read everything before believing any of it.

6 One engine, many sectors

The machine underneath is not about recruiting. It finds people matching a description, ranks them, writes to them, follows up, routes the replies, and keeps a record. Recruiting supplies the vocabulary, the data sources, and the compliance rules, not the mechanics. That distinction is the single most valuable structural decision in the product, and it is worth defending even when a shortcut would be faster.

What changes between sectors

- The words on screen: candidate or lead or subscriber, role or offer or launch, and whether a workspace reads as a client, a department, or a team.
- The criteria that matter: skills, seniority, and location for recruiting; company size, industry, tools in use, and buying signals for sales.
- The data sources: people databases and LinkedIn for recruiting; the same plus company and intent data for sales.
- The instructions the assistant writes under, since a recruiter's pitch and a seller's pitch are different jobs.
- The rules of the road: fair-hiring law for recruiting, marketing consent law for marketing.
- The default sequences, which differ in length, tone, and cadence by sector.

What does not change

- The funnel and every state in it.
- Approvals and the autonomy dial.
- Sending limits, quiet hours, warmup, and account health.
- The do-not-contact list and the audit record.
- Teams, workspace isolation, plans, and usage.
- The inbox, handoff, and reporting.

Put plainly, roughly nine tenths of what we are building is sector-neutral, and the last tenth is configuration. A second edition is not a second product.

The editions

Recruiter edition

NOW

Find candidates for an open role. The hardest compliance bar of the three, which is useful: building fair-hiring constraints, audit, and approval gates first means the later editions inherit a standard they would not otherwise have been held to.

Sales edition

NEXT

Find leads matching an ideal customer profile. Same funnel, different criteria and sources, plus a connection to the customer's CRM in place of the applicant tracking system. The nearest neighbour, and the cheapest to reach.

Nurture edition

LATER

Warm an existing list rather than a cold one, which changes where contacts come from and what consent means, but not what the engine does with them.

White label

The same structure makes a white-label offering straightforward. A partner puts their own name, logo, colours, domain, vocabulary, and default sequences on the product and sells it to their own customers, while the engine, the provider contracts, the compliance behaviour, and the operational limits stay ours and stay shared.

What makes this cheap is that settings resolve down a chain: our platform default, then the partner, then the customer's company, then the workspace, then the campaign, with the nearest value winning. A partner sets their vertical, their default sequences, and their sending limits once, and every customer they onboard inherits them. A customer can override company-wide, and a single workspace can tighten further. Nothing is forked, so adding a partner is data rather than a new build.

A partner sets defaults and carries the billing relationship. They do not get access to their customers' pipelines, and a partner is a label on a company rather than a layer above it, so no customer's data becomes reachable by virtue of who sold them the product.

There is a business reason to keep this option open beyond the extra revenue. Agencies and platforms bring distribution we would otherwise buy, and a partner who has put their own name on the product is far harder to displace than one reselling ours.

What this asks of us now

The option only stays open if we hold a few lines while building the recruiting edition. There is no separate "job" object, because a campaign already is one. Recruiting words live in a labels layer, never underneath. The audience description stays generic rather than growing recruiting-only fields. Data sources stay pluggable, so adding a sales database is adding an adapter. The assistant's sector-specific instructions stay in a swappable pack rather than in the main prompt. None of this costs meaningful time today, and all of it is expensive to retrofit.

FOCUS

The near-term goal is a recruiting product good enough that a regulated employer runs real roles through it, not a platform that does three things adequately. Everything above should be read as constraints on how we build the recruiting edition, not as work to schedule alongside it. We should be able to demonstrate a sales edition in a week when we choose to, and we should not choose to until recruiting is settled.

What "settled" means

- A recruiter can run a role end to end without leaving the product or asking us for help.
- Messages send from real connected mailboxes and LinkedIn seats, at real volume, without account trouble.
- A compliance reviewer can answer "why was this person contacted, and by whom" from the audit record alone.
- A team can move a campaign to Autopilot and keep it there for a month.
- Pipeline quality is good enough that recruiters accept most of what is proposed.

7 What we deliberately do not do

Each of these was considered and rejected for a reason worth keeping written down, because they come back up.

We do not send from our own LinkedIn accounts

One shared account messaging on behalf of many customers hits daily limits immediately, makes every message obviously automated, and is exactly the pattern LinkedIn takes to court. Each recruiter connects their own account, we operate it within human limits, and we stop at the first sign of trouble. This caps how fast LinkedIn can scale for us, which we accept.

We do not keep our own copy of a people database

Mirroring a provider's data would breach their terms, cost a great deal, go stale within months, and hand us a privacy liability covering millions of people we have no relationship with. We search live and keep only what a customer chooses to import.

We are not an applicant tracking system

The tracking system is the system of record for a hiring process, and customers already have one. We fill the top of the funnel and hand engaged candidates over. Competing with the record system would mean a much longer sale into a much more crowded market.

We do not store or archive customer communications for compliance

Regulated buyers in this space have not asked for it, and their existing mail archiving already captures it. We will build it if a deal turns on it.

Not in the first release

Deep integration with tracking systems, a browser extension, a mobile app, and the sales and nurture editions are all out of scope until the recruiting edition is settled. The first release is email-first, approval-first, and narrow on purpose.

8 Open decisions

These need a product call and are currently blocking or ambiguous.

- How a campaign decides it has enough candidates, and whether a recruiter sets a target number or the system infers one.
- How often a campaign should search: a simple fast, normal, or slow choice, or a number of hours.
- Whether a recruiter can tune what the fit score weighs, or whether we keep one scoring model everyone shares.
- What a "reject" does to a candidate: hide them from this campaign, from this workspace, or from everything.
- Whether templates are a company-level library that legal approves, or per-campaign text.
- Which tracking system we integrate with first, and whether handoff writes into it or only links to it.
- Whether Autopilot is available to every customer from the start, or unlocked after a track record.
- Whether the assistant's usage is metered to customers as a cost, and if so how we present it.
- Whether a partner can be given support access into a customer's workspace, and what the customer sees when they use it.
- Whether workspace-only do-not-contact entries are worth exposing at all, or whether every opt-out should stay company-wide.

9 Where we are today

Everything above describes the product as designed. This is an honest read of how much of it exists, so the document can be used for planning rather than only for alignment.

AREA	STATE	WHAT IS MISSING
Finding people	BUILT	One search strategy per pass rather than several; no automatic top-up target.
Ranking and fit	BUILT	No editable rubric; weights are the same for everyone.
Sequences and sending	BUILT	No template library; generated messages are not yet checked before queuing.
Inbox and replies	BUILT	Replies arriving by email do not yet reach the assistant.
Approvals	PARTLY	Approve and edit work; reject and snooze do not exist.
Connecting a mailbox	NOT STARTED	The screen exists but connecting a Gmail or Outlook account does nothing. This blocks the main channel.
Campaign control surface	PARTLY	No draft state or launch step, no per-campaign sources, channels, limits, targets, or threshold.
The assistant	PARTLY	Designs and sources well; never revisits a campaign, remembers nothing between runs, and chat cannot take actions.
Companies, workspaces, people	BUILT	Two levels, one login across companies, workspace grants, and person-owned seats are all in place. Partner branding is modelled but has no interface yet.
Compliance and safety	PARTLY	Audit and do-not-contact work; fair-hiring guarantees, unsubscribe wording in the message, data residency, and the account-wide kill switch do not exist.
Reporting	PARTLY	Funnel and reply rates work; per-sender health and message-level comparison do not.
Plans and usage	BUILT	Assistant usage is not counted, which is likely our largest variable cost.
Running it in production	NOT STARTED	No deployment setup, error tracking, or monitoring. Nothing is hosted anywhere yet.

The shape of the work is that the two engines, finding people and sending messages, are close to done, while the surface a customer configures and the safety story a compliance officer reads are the thin parts. Two items block a real pilot outright: connecting a work mailbox, and being able to run the product somewhere other than a laptop.

Working draft for internal discussion. Sections 1 to 7 describe intent and should be argued with; section 8 needs decisions; section 9 is a status read and will go out of date quickly.